

SUMMARY

Data engineering team lead who builds and runs data platforms end to end: a Databricks lakehouse, batch and streaming pipelines in Python, Spark and SQL, and dbt modeling with tests and data contracts. 10 years in data across logistics, retail and healthcare.

Currently leads a team of 4 data engineers at Harborlight Logistics, a freight-visibility SaaS, where the new lakehouse cut monthly warehouse cost by about 35%. Looking for a senior data platform role.

EXPERIENCE

Harborlight Logistics

Apr 2022 – Present

Data Engineering Team Lead (Jan 2024 – Present) · Senior Data Engineer (Apr 2022 – Dec 2023)

B2B SaaS for freight visibility used by shippers and carriers

- Lead a team of 4 data engineers that owns the data platform: carrier GPS/telematics and shipment-event ingestion, the lakehouse, dbt models and internal analytics
- Built the lakehouse on Databricks and Delta Lake, replacing Redshift and migrating 30 pipelines, which cut monthly warehouse cost by about 35%
- Built a near-real-time ETA events pipeline with Kinesis and Spark Structured Streaming, reducing latency from about 1 hour (batch) to under 5 minutes
- Introduced dbt with tests and set up data contracts with the backend team so schema changes stop breaking pipelines
- Work closely with data science, product and backend teams; write design docs and take part in the on-call rotation
- Interviewed and hired 2 engineers; mentor 2 junior engineers

Tech: Python, SQL, PySpark, Databricks, Delta Lake, dbt, Airflow (MWAA), AWS (S3, Kinesis, Lambda, Glue), Terraform, GitHub Actions, Docker

Data Engineer · Quillstone Retail

Jan 2019 – Mar 2022

- Built and maintained 40+ Airflow DAGs ingesting POS and e-commerce data into Amazon Redshift
- Processed clickstream data with Apache Spark on Amazon EMR for about a year
- Designed a dimensional data model for sales reporting, used by 120 analysts across 3 countries
- Migrated legacy SSIS jobs to Python, reducing nightly load time from 5 hours to 2 hours
- Optimized Redshift distribution and sort keys, cutting dashboard query times by 60%
- Contributed to data quality checks with Great Expectations

BI Developer · Pinecrest Health Analytics

Aug 2016 – Dec 2018

- Developed Tableau dashboards for hospital operations teams and monthly reporting for 14 client hospitals
- Wrote complex SQL for claims analysis in SQL Server

SKILLS

Languages: Python, SQL, Bash

Data platform: Databricks, Delta Lake, Apache Spark (PySpark, Structured Streaming), dbt, Airflow (MWAA), Amazon Redshift

Cloud & DevOps: AWS (S3, Kinesis, Lambda, Glue, EMR), Terraform, GitHub Actions, Docker

Modeling & quality: dimensional data modeling, dbt tests, Great Expectations, data contracts

BI: Tableau

EDUCATION

B.Sc. Industrial Engineering and Management · Ben-Gurion University of the Negev

2012 – 2016

MILITARY SERVICE

Operations Analyst · IDF

2009 – 2012

LANGUAGES

Hebrew (Native) · English (Fluent)